
Bandit Multiclass Classification

1 Problem Statement

A multiclass linear classifier is defined by a set of K vectors $w^{(1)}, \dots, w^{(K)} \in \mathbb{R}^d$ (or a compact matrix form $W = [w^{(1)} \ w^{(2)} \ \dots \ w^{(K)}]^\top \in \mathbb{R}^{K \times d}$), where K is the number of classes ($K \geq 3$) and d is the dimension of the feature vector. A set of samples $\{(x_t, y_t)\}_{t=1}^T$ where $x_t \in \mathbb{R}^d$ and $y_t \in [K]$ is called *linearly separable* if there exists a $w_*^{(1)}, \dots, w_*^{(K)}$ (or simply W_*) such that $y_t = \operatorname{argmax}_{y \in [K]} (W_* x_t)_y$ for all $t \in [T]$.

A set of samples $\{(x_t, y_t)\}_{t=1}^T$ is called *linearly separable with a margin γ* if there exists W_* with $\|w_*^{(i)}\|_2 \leq 1$ for all i and $(W_* x_t)_{y_t} \geq \max_{y \neq y_t} (W_* x_t)_y + \gamma$ for all t . We define the set $\mathcal{W} = \{W \in \mathbb{R}^{K \times d} : \|\mathbf{e}_i^\top W^*\|_2 \leq 1\}$. Therefore, $W_* \in \mathcal{W}$.

Consider the following *online multiclass classification* problem: at time t , the environment first reveals $x_t \in \mathbb{R}^d$ (and we assume $\|x_t\|_2 \leq 1$), then the learner predicts some label $\hat{y}_t \in [K]$; finally the environment reveals the true label $y_t \in [K]$. It is known that if the samples are linearly separable, then the learner will only make constant mistakes. It is achievable with the following simple perceptron algorithm:

Algorithm 1: Perceptron

```

1  $W_1 = \mathbf{0}$ 
2 for  $t = 1$  to  $T$  do
3   Predict  $\hat{y}_t = \operatorname{argmax}_y (W_t x_t)_y$ 
4   Update  $W_{t+1} \leftarrow W_t + (\mathbf{e}_{y_t} - \mathbf{e}_{\hat{y}_t}) x_t^\top$ 
5   (same as updating  $w_{t+1}^{(y_t)} \leftarrow w_t^{(y_t)} + x_t$  and  $w_{t+1}^{(\hat{y}_t)} \leftarrow w_t^{(\hat{y}_t)} - x_t$  when  $\hat{y}_t \neq y_t$ ).
```

The analysis is simple. On one hand we have (note $\langle A, B \rangle := \operatorname{Tr}(A^\top B)$):

$$\begin{aligned}
\langle W_t, W_* \rangle &= \langle W_{t-1}, W_* \rangle + \langle (\mathbf{e}_{y_t} - \mathbf{e}_{\hat{y}_t}) x_t^\top, W_* \rangle \\
&= \langle W_{t-1}, W_* \rangle + (\mathbf{e}_{y_t}^\top W_* x_t - \mathbf{e}_{\hat{y}_t}^\top W_* x_t) \\
&\geq \langle W_{t-1}, W_* \rangle + \gamma \mathbf{1}[\hat{y}_t \neq y_t].
\end{aligned}$$

By induction, we get $\langle W_{T+1}, W_* \rangle \geq \gamma \sum_{t=1}^T \mathbf{1}[\hat{y}_t \neq y_t]$.

On the other hand,

$$\langle W_{T+1}, W_* \rangle \leq \|W_{T+1}\|_F \|W_*\|_F \leq \sqrt{2K \sum_{t=1}^T \mathbf{1}[\hat{y}_t \neq y_t] \|x_t\|_2^2} \times 1 \leq \sqrt{2K \sum_{t=1}^T \mathbf{1}[\hat{y}_t \neq y_t]}$$

Combining the above two inequality, we get $\sum_{t=1}^T \mathbf{1}[\hat{y}_t \neq y_t] \leq \frac{4K}{\gamma^2}$.

The *bandit multiclass classification* problem proceeds in a very similar way. The only difference is that the learner only knows whether she predicts **correctly or not**, rather than gets the true label (i.e., in each round, the feedback is only $\mathbf{1}[\hat{y}_t = y_t]$, rather than y_t). Surprisingly, with this extremely

limited feedback, the learner can still make only constant mistakes. It can be achieved with the following variant of *halving algorithm*:

Algorithm 2: Halving

- 1 Discretize the space of W (i.e., the \mathcal{W} as defined in the beginning) with balls of radius $\frac{1}{2\gamma}$. Let the set of discretization points be \mathcal{S} . Then this is saying that for all $W \in \mathcal{W}$, there is always a $W' \in \mathcal{S}$ such that $\|W - W'\|_F \leq \frac{1}{2\gamma}$.
 - 2 Let $\mathcal{S}_1 = \mathcal{S}$
 - 3 **for** $t = 1$ **to** T **do**
 - 4 Let $\mathcal{S}_t(y) = \{W \in \mathcal{S}_t : (Wx_t)_y \geq (Wx_t)_i \forall i \in [K]\}$ (i.e., the set of W 's in \mathcal{S}_t that predict y as the label of x_t)
 - 5 Let $\tilde{y}_t = \operatorname{argmax}_y |\mathcal{S}_t(y)|$ (i.e., pick the majority vote)
 - 6 **if** $\tilde{y}_t \neq y_t$ **then**
 - 7 $\mathcal{S}_{t+1} = \mathcal{S}_t \setminus \mathcal{S}_t(\tilde{y}_t)$ (i.e., eliminate the W 's that are inconsistent with the outcomes)
 - 8 **else**
 - 9 $\mathcal{S}_{t+1} = \mathcal{S}_t$
-

Analysis. The majority in \mathcal{S}_t always has cardinality no less than $\frac{1}{K}|\mathcal{S}_t|$. So every time the learner predicts incorrectly, the size of $|\mathcal{S}_t|$ shrinks by an order of $(1 - \frac{1}{K})$. By our margin assumption, there is a $W' \in \mathcal{S}$ which incurs no error in all samples. Therefore, $|\mathcal{S}_t| \geq 1$ always holds. This means the number of errors is bounded by the order of $\frac{\log(|\mathcal{S}|)}{-\log(1 - \frac{1}{K})} \leq \mathcal{O}\left(K^2 d \log \frac{1}{\gamma}\right)$.

The main issue of this halving algorithm is that it is **inefficient** in the sense that $|\mathcal{S}|$ is in the order of $\frac{1}{\gamma^{Kd}}$.

Hence, we want to answer the following question: for bandit multiclass classification, can we have an efficient algorithm (time complexity polynomial in $K, d, \frac{1}{\gamma}, T$) that guarantees constant mistake bound (polynomial in $K, d, \frac{1}{\gamma}$) in the γ -separable case? (An unsolved open problem in [8])

2 Naive Approaches

- **Exhausting the feature space:** discretize the feature space into $\left(\frac{1}{\gamma}\right)^{\mathcal{O}(d)}$ blocks. The margin assumption guarantees that two points in a block would have the same label (so we can give every block a label). Each time a new point x_t comes, see whether the learner already knows the correct label for that block. If yes, predict that label; if not, randomly choose a label that has not been chosen for that block.
 \Rightarrow error bound: $K \left(\frac{1}{\gamma}\right)^{\mathcal{O}(d)}$.
- **Exhausting the hypothesis space:** Discretize the hypothesis space into $\left(\frac{1}{\gamma}\right)^{Kd}$ discretization points (like in the halving algorithm). For each of them, make predictions based on it until it makes an mistake.
 \Rightarrow error bound: $\left(\frac{1}{\gamma}\right)^{Kd}$.

3 Difficulties

Efficient classification algorithms often optimize a convex surrogate loss rather than a 0-1 loss. Examples include:

- Hinge loss: $\ell_t(W) = [1 - (Wx_t)_{y_t} + \max_{y \neq y_t} (Wx_t)_y]_+$.
- Logistic loss: $\ell_t(W) = -\ln \frac{\exp((Wx_t)_{y_t})}{\sum_y \exp((Wx_t)_y)}$
- A family of loss that interpolates hinge loss and squared hinge loss (see [2]'s Eq.(3))

In full-information setting, one can directly use online convex optimization techniques to deal with the above loss functions (learning over the space of W).

Many works for bandit classification reuse this kind of convex optimization schemes, together with explicit exploration and inverse propensity weighting [8, 7, 2, 6]. However, because it is hard to estimate these losses when the true label y_t is not known, some of these algorithms [2, 6] simply do not update when the learner makes a mistake ($\tilde{y}_t \neq y_t$). We give this kind of algorithm an error lower bound $\Omega\left(\left(\frac{1}{\gamma}\right)^{(d-1)/2}\right)$ in Section 10.

It indeed looks hard to design a convex loss for W 's when the learner makes a mistake: when $\tilde{y}_t \neq y_t$, the set of W 's that we want to penalize (i.e., to assign larger loss) is $\{W : (Wx_t)_{\tilde{y}_t} \geq (Wx_t)_y, \forall y\}$, which is a convex cone in the space of W . It is impossible for a convex function to be large only in a convex subset (the case $K = 2$, on the other hand, does not have this issue). Can we argue that if the learner is restricted to use convex losses in the space of W , she will have to suffer exponential (in K) errors?

4 One-versus-all Perceptron

Our assumption of linearly separable with a margin is

$$(W_*x_t)_{y_t} \geq (W_*x_t)_y + \gamma, \forall y \neq y_t. \quad (1)$$

As discussed in Section 3, it seems hard to make update when $\tilde{y}_t \neq y_t$ and achieve a constant and polynomial error bound.

For a moment, in this section we make the following stronger margin assumption (which we call *one-versus-all separable* with a margin):

$$\begin{cases} (W_*x_t)_{y_t} \geq \gamma/2 \\ (W_*x_t)_y \leq -\gamma/2, \forall y \neq y_t. \end{cases} \quad (2)$$

Clearly, one-versus-all separability implies linearly separability, but not the other way around. With one-versus-all separability assumption, we can view the problem as K parallel binary classification problem. The following algorithm achieves constant error bound:

Algorithm 3: One-versus-all Perceptron

```

1 Initialize  $w_t^{(1)} = \dots = w_t^{(K)} = \mathbf{0} \in \mathbb{R}^d$ 
2 for  $t = 1$  to  $T$  do
3   if  $\exists y$  such that  $w_t^{(y)\top} x_t \geq 0$  then
4     Assign  $\tilde{y}_t$  to any  $y$  with  $w_t^{(y)\top} x_t \geq 0$ 
5     Predict  $\tilde{y}_t$ 
6     if  $\tilde{y}_t \neq y_t$  then update  $w_t^{(\tilde{y}_t)} \leftarrow w_t^{(\tilde{y}_t)} - x_t$ ;
7   else
8     Pick  $\tilde{y}_t$  randomly from  $\text{unif}[K]$ 
9     Predict  $\tilde{y}_t$ 
10    if  $\tilde{y}_t = y_t$  then update  $w_t^{(\tilde{y}_t)} \leftarrow w_t^{(\tilde{y}_t)} + x_t$ ;

```

Analysis. Note that when the algorithm enters Line 6 or Line 10, the binary classifier $w_t^{(\tilde{y}_t)}$ is making an error. Let the number of times the algorithm enters Line 6 and Line 10 be M and N respectively. By the error bound of binary perceptron $\mathcal{O}\left(\frac{1}{\gamma^2}\right)$, we have that $M + N = \mathcal{O}\left(\frac{K}{\gamma^2}\right)$. Then note that the number of times the algorithm makes a mistake (i.e. $\tilde{y}_t \neq y_t$) is upper bounded by M plus how many times the algorithm explores in Line 8; and when the algorithm explores, with probability $\frac{1}{K}$ it enters Line 10. Therefore, the number of mistakes is bounded (in expectation) by $\mathbb{E}[M + KN] = \mathcal{O}\left(\frac{K^2}{\gamma^2}\right)$.

5 One-versus-all Kernel Perceptron

The idea here is to do a feature transformation such that if the original assumption (1) is satisfied, then in the transformed feature space, (2) will be satisfied, then we can run the one-versus-all perceptron given in 4. Under the original assumption, each class corresponds to an intersection of $K - 1$ halfspaces in the feature spaces (i.e., class i corresponds to $\{x \in \mathbb{R}^d : w_*^{(i)\top} x \geq w_*^{(j)\top} x, \forall j \neq i\}$). We want this region to become a halfspace in the transformed space. Fortunately, this is possible when there is a margin, and is exactly what is done in [9].

5.1 Refinement for [9]’s Fact 1

Lemma 1. For $i = 1, \dots, \ell$ let $q^{(i)}(x) = \sum_S c_S^{(i)} x_S$ be a polynomial over x_1, \dots, x_d with $\|q^{(i)}\|^2 \leq M_i$. Then (1) if $q^{(1)} \dots q^{(\ell)}$ has degree at most \deg , we have $\|q^{(1)} \dots q^{(\ell)}\|^2 \leq \ell^{\deg} \prod_i M_i$, and (2) we have $\|q^{(1)} + \dots + q^{(\ell)}\|^2 \leq \ell(M_1 + \dots + M_\ell)$.

Proof. For the first bound, we bound the ratio between the following two values: $\|q^{(1)} \dots q^{(\ell)}\|^2$ and $\|q^{(1)}\|^2 \dots \|q^{(\ell)}\|^2$.

$$\begin{aligned} n_1 &= \prod_{i=1}^{\ell} \|q^{(i)}\|^2 = \prod_{i=1}^{\ell} \left(\sum_{S_i} \left(c_{S_i}^{(i)} \right)^2 \right) = \sum_{(S_1, \dots, S_\ell)} \left(\prod_{i=1}^{\ell} \left(c_{S_i}^{(i)} \right)^2 \right) = \sum_{(S_1, \dots, S_\ell)} \left(\prod_{i=1}^{\ell} c_{S_i}^{(i)} \right)^2 \\ &= \sum_S \sum_{(S_1, \dots, S_\ell): S_1 \times \dots \times S_\ell = S} \left(\prod_{i=1}^{\ell} c_{S_i}^{(i)} \right)^2 \end{aligned} \quad (3)$$

$$n_2 = \|q^{(1)} \dots q^{(\ell)}\|^2 = \sum_S c_S^2 = \sum_S \left(\sum_{(S_1, \dots, S_\ell): S_1 \times \dots \times S_\ell = S} \left(c_{S_1}^{(1)} \dots c_{S_\ell}^{(\ell)} \right) \right)^2 = \sum_S \left(\sum_{(S_1, \dots, S_\ell): S_1 \times \dots \times S_\ell = S} \prod_{i=1}^{\ell} c_{S_i}^{(i)} \right)^2 \quad (4)$$

Let M be an upper bound of the number of terms involved in the summation $\sum_{(S_1, \dots, S_\ell): S_1 \times \dots \times S_\ell = S}$, then by Cauchy-Schwarz’s inequality we have $n_2 \leq M n_1$.

M counts the number of different (S_1, \dots, S_ℓ) ’s with $S_1 \times \dots \times S_\ell = S$. Since S has degree at most \deg , we can bound M by ℓ^{\deg} .

The second bound can be obtained by applying Cauchy-Schwarz once. □

5.2 Refinement for [9]’s Theorem 10

Lemma 2. Under the same condition as in [9]’s Theorem 10, the PTF has margin on X is at least $(1/t)^{O(r \log r + r \log \log t)}$.

Proof. We simply follow the construction in the original paper. $\|1 - w^i \cdot x\|^2 \leq 4$ can imply $\|(1 - w^i \cdot x)^j\|^2 \leq j^{O(j)} 4^j \leq j^{O(j)}$. Then $\|a_j(1 - w^i \cdot x)^j\|^2 \leq 2^{2r} r^{O(r)} = r^{O(r)}$ follows. Thus $\|T_r(1 - w^i \cdot x)\|^2 = \|\sum_{j=0}^r a_j(1 - w^i \cdot x)^j\|^2 \leq (r+1)^2 r^{O(r)} = r^{O(r)}$. Then, $(P(w^i \cdot x))^{\log 2t} \leq (\log t)^{O(r \log t)} \times r^{O(r \log t)} \leq (r \log t)^{O(r \log t)}$. Finally, $\|p\|^2 \leq (t+1)^2 (r \log t)^{O(r \log t)} \leq (t+1)^2 (t^{O(r \log r + r \log \log t)}) = t^{O(r \log r + r \log \log t)}$. □

5.3 Refinement for [9]’s Theorem 7

Lemma 3. Under the same condition as in [9]’s Theorem 7, the PTF has margin on X is at least $\left(\frac{\rho}{t \log t \log \frac{1}{\rho}} \right)^{O(t \log t \log 1/\rho)}$.

Proof. Using the same construction, we bound $\|p\|$. First, we have $\|\frac{2w^i \cdot x}{\rho}\|^2 \leq \frac{4}{\rho^2}$ and $\|\left(\frac{2w^i \cdot x}{\rho}\right)^j\|^2 \leq j^j \left(\frac{4}{\rho^2}\right)^j = \left(\frac{4j}{\rho^2}\right)^j$. $a(x), b(x)$ are polynomials of degree $O(\log t \log \frac{1}{\rho})$ with coefficients of magnitude $\left(\frac{1}{\rho}\right)^{O(\log t \log 1/\rho)}$. Thus

$$\begin{aligned} \|a(2w^i \cdot x/\rho)\|^2 &\leq \left(\frac{\log t \log \frac{1}{\rho}}{\rho}\right)^{O(\log t \log 1/\rho)} \times \left(\frac{1}{\rho}\right)^{O(\log t \log 1/\rho)} \\ &\leq \left(\frac{\log t \log \frac{1}{\rho}}{\rho}\right)^{O(\log t \log 1/\rho)}. \end{aligned}$$

Same holds for $\|b(2w^i \cdot x/\rho)\|^2$. Finally we have $\|B(x)\|^2 \leq (t)^{O(\log t \log 1/\rho)} \times \left(\frac{\log t \log \frac{1}{\rho}}{\rho}\right)^{O(\log t \log 1/\rho)} = \left(\frac{t \log t \log \frac{1}{\rho}}{\rho}\right)^{O(\log t \log 1/\rho)}$

[TODO] Refined lower bound for $|B(x)|$

□

6 A Recipe for Designing Efficient Algorithms

We have seen that if the data is “one-versus-all” separable with a margin, then constant error is achievable. To have constant error, a key is that every time there is a mistake, the learner should make significant progress.

Besides one-versus-all separability, are there other kinds of conditions under which we can easily make progress when we make a mistake? If such condition exists, then we have the following recipe of designing algorithms that have constant mistake bound:

- Design an algorithm that has constant mistake bound under some assumption.
- Design feature transformation such that our original linear separability satisfies the above assumption in the transformed space.

Our kernel perceptron exactly follows the above procedure. Following is another example.

6.1 One-sided banditron

Consider the following algorithm.

Algorithm 4: Ond-sided banditron

- 1 **Define:** $\Omega \subseteq \mathbb{R}^{K \times d}$ is a convex set, and we assume $W^* \in \Omega_t$ for all t . Learning rate η .
 - 2 **Initialize:** $W_1 = 0$.
 - 3 **for** $t = 1, \dots, T$ **do**
 - 4 Receive $x_t \in \mathbb{R}^d$.
 - 5 Predict $\hat{y}_t = \operatorname{argmax}_{y \in [K]} (W_t x_t)_y$.
 - 6 **if** $\hat{y}_t \neq y_t$ **then**
 - 7 $W'_{t+1} \leftarrow W_t - \eta \mathbf{e}_{\hat{y}_t} x_t^\top$.
 - 8 $W_{t+1} = \Pi_{\Omega_{t+1}}(W'_{t+1}) \triangleq \operatorname{argmin}_{W \in \Omega_{t+1}} \|W - W'_{t+1}\|_F$.
-

The above algorithm is similar to banditron but only performs *one-sided* update and never explicitly explores to get label. Below we check under what condition can this algorithm have constant mistake bound.

Following standard analysis for online gradient descent, we have when the learner makes a mistake:

$$\begin{aligned} \|W_{t+1} - W^*\|_F^2 &\leq \|W'_{t+1} - W^*\|_F^2 = \|W_t - W^* - \eta \mathbf{e}_{\hat{y}_t} x_t^\top\|_F^2 \\ &= \|W_t - W^*\|_F^2 - 2\eta((W_t x_t)_{\hat{y}_t} - (W^* x_t)_{\hat{y}_t}) + \eta^2 \|x_t\|_2^2 \end{aligned} \quad (5)$$

We see that if the condition

$$(W_t x_t)_{\hat{y}_t} - (W^* x_t)_{\hat{y}_t} \geq \Delta \quad (6)$$

holds whenever $\hat{y}_t \neq y_t$ and we pick $\eta = \Delta / \sup_{\tau} \|x_{\tau}\|_2^2$, then we always have

$$\|W_{t+1} - W^*\|_F^2 \leq \|W_t - W^*\|_F^2 - \frac{2\Delta^2}{\sup_{\tau} \|x_{\tau}\|_2^2} + \frac{\Delta^2}{\sup_{\tau} \|x_{\tau}\|_2^2} = \|W_t - W^*\|_F^2 - \frac{\Delta^2}{\sup_{\tau} \|x_{\tau}\|_2^2}$$

whenever $\hat{y}_t \neq y_t$. Then we have the error bound

$$\sum_{t=1}^T \mathbf{1}[\hat{y}_t \neq y_t] \leq \frac{\|W^*\|_F^2 \sup_{\tau} \|x_{\tau}\|_2^2}{\Delta^2}$$

because $\|W_t - W^*\|_F^2$ is always non-negative and $\|W_1 - W^*\|_F^2 = \|W^*\|_F^2$.

The following analysis leads to the same bound but is more standard. Let $M_t = \mathbf{1}[\hat{y}_t \neq y_t]$. By (5) and (6) we have

$$\Delta \sum_{t=1}^T M_t \leq \sum_{t=1}^T M_t ((W_t x_t)_{\hat{y}_t} - (W^* x_t)_{\hat{y}_t}) \leq \frac{\|W^*\|_F^2}{2\eta} + \frac{\eta}{2} \sup_{\tau} \|x_{\tau}\|_2^2 \sum_{t=1}^T M_t$$

Picking appropriate η and solving for $\sum_{t=1}^T M_t$ we get the same error bound.

A candidate set of assumptions that leads to (6) is the *sum-to-zero + margin* assumption:

- $\sum_{i=1}^K \mathbf{e}_i^{\top} W^* = 0$ (so we pick $\Omega_t = \{W \in \mathbb{R}^{K \times d} : \|\mathbf{e}_i^{\top} W\|_2 \leq 1, \sum_{i=1}^K \mathbf{e}_i^{\top} W = 0\}$)
- $(W^* x_t)_y \leq -\Delta$ for all $y \neq y_t$.

They are similar to the assumptions made in [5]. We see that the above two conditions indeed lead to (6): $(W_t x_t)_{\hat{y}_t} = \max_y (W_t x_t)_y \geq \frac{1}{K} \sum_{i=1}^K (W_t x_t)_i = 0$ (recall that we always project W_t back to Ω_t), and when $\hat{y}_t \neq y_t$ we have $(W^* x_t)_{\hat{y}_t} \leq -\Delta$. Unfortunately, this set of assumptions is even stronger than the one-versus-all separable assumption.

Although the one-versus-all assumption is similar, there we do not perform projection for W_t , so it is possible that $\max_y (W_t x_t)_y < 0$. In that case, what we do is uniformly sample a label.

Questions:

- Can we avoid the projection step for one-sided banditron (just like perceptron)?
- Are there other natural assumptions that guarantee (6)?

6.2 Algorithms that can update without true labels

Algorithm 4 can update without the information of y_t . One can think of Algorithm 4 as a gradient algorithm on the loss function $\ell_t(W) = \mathbf{1}[\hat{y}_t \neq y_t](W x_t)_{\hat{y}_t}$. As a comparison, one can think of the surrogate loss function in [5] as $\ell_t(W) = \sum_{y=1}^K \mathbf{1}[y \neq y_t][\Delta + (W x_t)_y]_+$ (some scaling and notations are changed for better comparison), whose inverse propensity weighting estimator becomes $\tilde{\ell}_t(W) = \frac{1}{p_{t,y_t}} \mathbf{1}[\hat{y}_t \neq y_t][\Delta + (W x_t)_{\hat{y}_t}]_+$. One can see that they are closely related, except that in [5] when $(W x_t)_{\hat{y}_t} \leq -\Delta$, the loss for W is set to zero (which makes no difference under the sum-to-zero + margin assumption).

7 Logarithmic Regret for Strongly Convex and Smooth Losses

7.1 Gradient descent

In some previous works (we will visit them in the next subsection), it is assumed that the labels are generated with the following process:

$$p(y|x) = \nabla_y \Phi(W^* x), \quad (7)$$

with some $\Phi : \mathbb{R}^K \rightarrow \mathbb{R}$. This is a *generalized linear model*. The margin assumption is defined as

$$p(y_t|x_t) - \max_{y \neq y_t} p(y|x_t) \geq \Delta, \forall t. \quad (8)$$

This margin assumption is a special case of the Tsybakov condition investigated in [1]. Defining $\hat{y}_t = \operatorname{argmax}_y \nabla \Phi(W_t x_t)$ and $y_t^* = \operatorname{argmax}_y \nabla \Phi(W^* x_t)$ The regret is defined as

$$\begin{aligned} \operatorname{Reg}_T &= \sum_{t=1}^T \mathbb{E} [\mathbf{1}[\hat{y}_t \neq y_t] - \mathbf{1}[y_t^* \neq y_t]] \\ &\text{(the } \mathbb{E}[\cdot] \text{ is conditioned on all history before time } t \text{ and the randomness is from } y_t) \\ &= \sum_{t=1}^T \sum_{i=1}^K (\mathbf{1}[\hat{y}_t \neq i] - \mathbf{1}[y_t^* \neq i]) \nabla_i \Phi(W^* x_t). \end{aligned} \quad (9)$$

Below we argue that if Φ is strongly convex and smooth, then the following algorithm (almost exactly the algorithm in [1] except that we replace label query by uniform exploration) achieves regret logarithmic in T .

Assumption 1. For any u, v ,

$$\Phi(v) + \langle \nabla \Phi(v), u - v \rangle + \frac{\gamma_\ell}{2} \|u - v\|_2^2 \leq \Phi(u) \leq \Phi(v) + \langle \nabla \Phi(v), u - v \rangle + \frac{\gamma_u}{2} \|u - v\|_2^2.$$

That is, Φ is γ_ℓ -strongly convex and γ_u -smooth.

Algorithm 5: Selectron [1]

```

1 Definition:  $\epsilon \triangleq \Delta / (2\gamma_u C')$ , where  $C'$  is defined in Lemma 6,
2  $\eta \triangleq \frac{\sup_{W \in \Omega} \|W\|_F^2}{\gamma_\ell}$ ,
3  $\Omega \triangleq$  a convex set where  $W^*$  lies in.
4  $\ell_t(W) = \Phi(W x_t) - (W x_t)_{y_t}$ .
5 Initialization:  $W_1 = 0, M_1 = I$ .
6 for  $t = 1, \dots, T$  do
7   Observe  $x_t$ .
8   if  $\|x_t\|_{M_t^{-1}} \geq \epsilon$  then
9     Draw  $\tilde{y}_t \sim \text{unif}([K])$ .
10  else
11    Draw  $\tilde{y}_t = \hat{y}_t \triangleq \operatorname{argmax}_{y \in [K]} (W_t x_t)_y$ .
12  if  $\tilde{y}_t = y_t$  then
13     $Z_t \leftarrow 1$ ,
14     $M_{t+1} \leftarrow M_t + Z_t x_t x_t^\top$ ,
15     $W_{t+1} = \operatorname{argmin}_{W \in \Omega} \left\{ \sum_{s=1}^t Z_s \ell_s(W) + \frac{1}{2\eta} \|W\|_F^2 \right\}$ .
16  else
17     $Z_t \leftarrow 0$ ,
18     $M_{t+1} \leftarrow M_t$ ,
19     $W_{t+1} \leftarrow W_t$ .

```

Theorem 4. The selective sampling algorithm achieves

$$\operatorname{Reg}_T \leq \tilde{O} \left(\frac{\gamma_u^2 K^2 d}{\gamma_\ell^2 \Delta^2} \right).$$

The proof of this theorem is decomposed to the following three lemmas.

Definition 5. $\|W\|_M^2 \triangleq \sum_{i=1}^K \|\mathbf{e}_i^\top W\|_M^2$.

With this definition we have $\|W x_t\|_2^2 = \sum_{i=1}^K (\mathbf{e}_i^\top W x_t)^2 \leq \sum_{i=1}^K \|\mathbf{e}_i^\top W\|_M^2 \|x_t\|_{M^{-1}}^2 \leq \|W\|_M^2 \|x_t\|_{M^{-1}}^2$

Lemma 6 (Proposition 1 of [1]). *With probability at least $1 - \delta$,*

$$\|W_t - W^*\|_{M_t} \leq C' \triangleq \frac{2C\sqrt{dK}}{\gamma_\ell} + \sqrt{\frac{2}{\eta\gamma_\ell}} \sup_{W \in \Omega} \|W\|_F,$$

where C hides some $\log t$ and $\log(1/\delta)$ terms (defined in the proof).

Proof. By the first-order optimality condition on W_t we have

$$0 \geq \left\langle \frac{1}{\eta} W_t + \sum_{s=1}^{t-1} Z_s \nabla \ell_s(W_t), W_t - W^* \right\rangle.$$

By the definition of ℓ_t ,

$$\begin{aligned} 0 &\geq \left\langle \frac{1}{\eta} W_t + \sum_{s=1}^{t-1} Z_s \nabla \ell_s(W_t), W_t - W^* \right\rangle \\ &= \left\langle \sum_{s=1}^{t-1} Z_s (\nabla \Phi(W_t x_s) - \mathbf{e}_{y_s}) x_s^\top, W_t - W^* \right\rangle + \frac{1}{\eta} \langle W_t, W_t - W^* \rangle \\ &= \left\langle \sum_{s=1}^{t-1} Z_s (\nabla \Phi(W_t x_s) - \nabla \Phi(W^* x_s) + \xi_s) x_s^\top, W_t - W^* \right\rangle + \frac{1}{\eta} \langle W_t, W_t - W^* \rangle, \\ &\quad (\text{let } \xi_s \triangleq \nabla \Phi(W^* x_s) - \mathbf{e}_{y_s}; \text{ note that } \mathbb{E}[\xi_s] = \mathbf{0}) \\ &= \left\langle \sum_{s=1}^{t-1} Z_s (\nabla \Phi(W_t x_s) - \nabla \Phi(W^* x_s) + \xi_s), W_t x_s - W^* x_s \right\rangle + \frac{1}{\eta} \langle W_t, W_t - W^* \rangle, \\ &\geq \gamma_\ell \sum_{s=1}^{t-1} Z_s \|W_t x_s - W^* x_s\|_2^2 + \sum_{s=1}^{t-1} Z_s \langle \xi_s, W_t x_s - W^* x_s \rangle + \frac{1}{\eta} \langle W_t, W_t - W^* \rangle. \\ &\quad (\text{for } \gamma_\ell\text{-strongly convex functions, } \langle \nabla \Phi(u) - \nabla \Phi(v), u - v \rangle \geq \gamma_\ell \|u - v\|_2^2) \end{aligned}$$

Rearranging gives

$$\begin{aligned} \sum_{s=1}^{t-1} Z_s \langle \xi_s, W^* x_s - W_t x_s \rangle &\geq \gamma_\ell \sum_{s=1}^{t-1} Z_s \|W_t x_s - W^* x_s\|_2^2 - \frac{2}{\eta} \sup_{W \in \Omega} \|W\|_F^2 \\ &= \gamma_\ell \sum_{s=1}^{t-1} Z_s \sum_{i=1}^K \langle \mathbf{e}_i W_t - \mathbf{e}_i W^*, x_s \rangle^2 - \frac{2}{\eta} \sup_{W \in \Omega} \|W\|_F^2 \\ &= \gamma_\ell \|W_t - W^*\|_{M_t}^2 - \frac{2}{\eta} \sup_{W \in \Omega} \|W\|_F^2. \end{aligned} \tag{10}$$

On the other hand,

$$\begin{aligned} \sum_{s=1}^{t-1} Z_s \langle \xi_s, W^* x_s - W_t x_s \rangle &= \sum_{i=1}^K \sum_{s=1}^{t-1} Z_s \xi_{s,i} ((W^* x_s)_i - (W_t x_s)_i) \\ &= \sum_{i=1}^K \sum_{s=1}^{t-1} Z_s \xi_{s,i} \langle W_i^* - W_{t,i}, x_s \rangle \\ &\leq \sum_{i=1}^K \left\| \sum_{s=1}^{t-1} Z_s \xi_{s,i} x_s \right\|_{M_t^{-1}} \|W_i^* - W_{t,i}\|_{M_t}. \end{aligned} \tag{11}$$

Note that $\{\xi_{s,i}\}_s$ is a martingale difference sequence. Using [4]'s Lemma 1 in Appendix A.1 with constant $c_m = \sup_\tau \|x_\tau\|_2 = 1$, $\lambda_0 = \frac{1}{\eta\gamma_\ell}$, $R = 2$ yields with probability at least $1 - \delta/K$

$$\left\| \sum_{s=1}^{t-1} Z_s \xi_{s,i} x_s \right\|_{M_t^{-1}} \leq 2\sqrt{3 + 2\log(1 + \eta\gamma_\ell)} \sqrt{d \log t} \sqrt{\log(dK/\delta)} \triangleq \sqrt{d}C.$$

Thus, by (11) we get with probability $1 - \delta$

$$\sum_{s=1}^{t-1} Z_s \langle \xi_s, W^* x_s - W_t x_s \rangle = C\sqrt{d} \sum_{i=1}^K \|W_i^* - W_{t,i}\|_{M_t} = C\sqrt{dK} \|W^* - W_t\|_{M_t},$$

where in the last inequality we use $\sum_{i=1}^K a_i \leq \sqrt{K} \sqrt{\sum_{i=1}^K a_i^2}$. Combining with (10) and solving for $\|W_t - W^*\|_{M_t}$ we get

$$\|W_t - W^*\|_{M_t} \leq \frac{2C\sqrt{dK}}{\gamma_\ell} + \sqrt{\frac{2}{\eta\gamma_\ell}} \sup_{W \in \Omega} \|W\|_F.$$

□

Lemma 7. *If $\|x_t\|_{M_t^{-1}} \leq \epsilon = \Delta/(2\gamma_u C')$, then $\hat{y}_t = y_t^*$.*

Proof.

$$\begin{aligned} \nabla \Phi_{y_t^*}(W^* x_t) - \nabla \Phi_{\hat{y}_t}(W^* x_t) &\leq \nabla \Phi_{y_t^*}(W^* x_t) - \nabla \Phi_{\hat{y}_t}(W^* x_t) - \nabla \Phi_{y_t^*}(W_t x_t) + \nabla \Phi_{\hat{y}_t}(W_t x_t) \\ &\leq 2\gamma_u \|W^* x_t - W_t x_t\|_2 \\ &\leq 2\gamma_u \|W^* - W_t\|_{M_t} \|x_t\|_{M_t^{-1}} \\ &\leq \Delta. \end{aligned} \quad (\text{by Lemma 6})$$

By the margin assumption (8), $\hat{y}_t = y_t^*$. □

From Lemma 7 and the regret definition (9), we see that in the rounds $\|x_t\|_{M_t^{-1}} \leq \epsilon$, the regret is zero. The following lemma bounds the number of rounds with $\|x_t\|_{M_t^{-1}} > \epsilon$

Lemma 8. $\sum_{t=1}^T \mathbf{1}[\|x_t\|_{M_t^{-1}} > \epsilon] \leq \frac{K \ln T \ln \frac{1}{\delta}}{\epsilon^2}$ with probability at least $1 - \delta$.

Proof.

$$\begin{aligned} &\sum_{t=1}^T \mathbf{1}[\|x_t\|_{M_t^{-1}} > \epsilon] \\ &\leq \left(K \ln \frac{1}{\delta} \right) \sum_{t=1}^T \mathbf{1}[\|x_t\|_{M_t^{-1}} > \epsilon] Z_t \quad (\text{when } \|x_t\|_{M_t^{-1}} \geq \epsilon, \hat{y}_t = y_t \text{ with probability } \frac{1}{K}) \\ &\leq \frac{K \ln \frac{1}{\delta}}{\epsilon^2} \sum_{t=1}^T \|x_t\|_{M_t^{-1}}^2 Z_t \leq \frac{K \ln T \ln \frac{1}{\delta}}{\epsilon^2}. \end{aligned}$$

□

Proof of Theorem 4. By Lemma 7 and 8,

$$\begin{aligned} \text{Reg}_T &= \sum_{t=1}^T \sum_{i=1}^K (\mathbf{1}[\hat{y}_t \neq i] - \mathbf{1}[y_t^* \neq i]) \nabla_i \Phi(W^* x_t) \\ &\leq \sum_{t=1}^T \left(\mathbf{1}[\|x_t\|_{M_t^{-1}} > \epsilon] \right) \sum_{i=1}^K \nabla_i \Phi(W^* x_t) \\ &\leq \sum_{t=1}^T \left(\mathbf{1}[\|x_t\|_{M_t^{-1}} > \epsilon] \right) \\ &\leq \tilde{O} \left(\frac{K \gamma_u^2 C'^2}{\Delta^2} \right). \end{aligned}$$

Using the definition of C' we get the final bound. □

7.2 Revisiting the stochastic assumptions made in [3] and [1]

In [3] and [1], the feature vectors x_t 's are adversarially picked, while y_t is generated based on some conditional distribution based on x_t . For example, in [3], the label is generated by

$$p(y_t = i | x_t) = \frac{\alpha + (W^* x_t)_i}{\alpha + 1}, \quad (12)$$

where $W^* \in \Omega_t \triangleq \{W \in \mathbb{R}^{K \times d} : \sum_{i=1}^K (W x_t)_i = \alpha + 1 - K\alpha \text{ and } -\alpha \leq (W x_t)_i\}$; a valid Φ that satisfies (8) is

$$\Phi(u) = \frac{1}{2(\alpha + 1)} \|u\|_2^2 + \frac{\alpha}{\alpha + 1} \sum_{i=1}^K u_i. \quad (13)$$

in [1], the labels are generated by

$$p(y_t = i | x_t) = \frac{\exp((W^* x_t)_i)}{\sum_{j=1}^K \exp((W^* x_t)_j)}, \quad (14)$$

where $W^* \in \Omega \triangleq \{W \in \mathbb{R}^{K \times d} : \|\mathbf{e}_i^\top W\|_2 \leq D\}$, and Φ can be chosen as

$$\Phi(u) = \log \left(\sum_{i=1}^K \exp(u_i) \right). \quad (15)$$

If we want, we can also cast our original linear separability to this framework. But it seems we need to incorporate the margin assumption (8) into the definition of Φ :

$$p(y_t = i | x_t) = \mathbf{1}[(W x_t)_i \geq (W x_t)_j + \Delta, \forall j \neq i]. \quad (16)$$

A natural definition of Φ that incorporates the margin assumption (8) is

$$\Phi(u) = \sum_i u_i \mathbf{1}[u_i \geq u_j + \Delta, \forall j \neq i] \quad (17)$$

Discussions:

- Although [7, 6] do not make stochastic assumptions, their loss functions are indeed strongly-convex and smooth (although γ_u/γ_ℓ can be quite large). It is worth notice that [6] also obtained some logarithmic bound (their Theorem 3), but the bound is margin-independent. It seems that their Theorem 3 is tighter than our bound here. Not sure if it is because they use improper learning (we indeed should see whether improper learning helps in our problem). There is a α parameter in [7], somehow it looks related to margin.
- How is the assumptions/algorithms/bounds introduced in this section related to [10]? They are similar in that they both assume $p(y|x)$ to be defined by some function in a known function class. The margin assumption (8) is a generalization of [10]'s to the multiclass case. One can also view the FTRL step used in Selectron as the counterpart of the ridge regression step defined in [10].
- We indeed have tried to use this kind of selective sampling idea to get finite error bound (see Section 8). But it seems it is impossible/difficult to find a $\Phi(\cdot)$ that at the same time 1) is strongly convex, and 2) whose logarithmic/constant regret bound implies logarithmic/constant mistake bound.

8 A KWIK Algorithm [TODO]

Assumption 2. $\|x_t\|_2^2 \leq 1$. There is a $W^* \in \mathcal{W}$ such that $\ell_t(W^*) \leq 0$ for all t (ℓ_t and \mathcal{W} are defined below).

Algorithm 6: KWIK Banditron

1 **Input:** $D \geq 2, \epsilon$ (picked in a later lemma).

2 **Definition:**

$$\begin{aligned}\ell_t(W) &\triangleq [1 - (Wx_t)_{y_t} + \max_{r \neq y_t} (Wx_t)_r]_+^2 \quad (\text{squared hinge loss}) \\ &= \Phi_t(Wx_t),\end{aligned}$$

where $\Phi_t(z) \triangleq [1 - \mathbf{e}_{y_t}^\top z + \max_{r \neq y_t} \mathbf{e}_r^\top z]_+^2$.

3 Also, define $\mathcal{W} = \{W \in \mathbb{R}^{K \times d} : \|\mathbf{e}_i^\top W\|_2 \leq D \text{ for all } i \in [K]\}$.

4 **Initialization:** $W_1 = 0, M_1 = I$.

5 **for** $t = 1, \dots, T$ **do**

6 Observe x_t .

7 **if** $\|x_t\|_{M_t^{-1}} \geq \epsilon$ **and** $\|W_t - W^*\|_F \geq 1$ **then**

8 Draw $\tilde{y}_t \sim \text{unif}([K])$.

9 **else**

10 Draw $\tilde{y}_t = \hat{y}_t \triangleq \arg\max_{r \in [K]} (W_t x_t)_r$.

11 **if** $\tilde{y}_t = y_t$ **then**

12 $Z_t \leftarrow 1$,

13 $M_{t+1} \leftarrow M_t + Z_t \ell_t(W_t) x_t x_t^\top$,

14 $W_{t+1} \leftarrow \Pi_{\mathcal{W}}(W_t - \eta_{t+1} \nabla \ell_t(W_t))$, where $\eta_{t+1} = \frac{1}{8}$.

15 ($\Pi_{\mathcal{W}}$ is the projection operator onto \mathcal{W} w.r.t. Frobenius norm)

16 **else**

17 $Z_t \leftarrow 0$,

18 $M_{t+1} \leftarrow M_t$,

19 $W_{t+1} \leftarrow W_t$.

Lemma 9. $\|\nabla \ell_t(W)\|_F^2 \leq 8\ell_t(W)$.

Proof. $\|\nabla \ell_t(W)\|_F^2 = \|\nabla \Phi_t(Wx_t)x_t^\top\|_F^2 \leq \left(2\sqrt{\Phi_t(Wx_t)}\right)^2 \times 2\|x_t\|_2^2 \leq 8\ell_t(W)$. \square

Lemma 10. Let $L_{t+1} \triangleq \sum_{s=1}^t Z_s \ell_s(W_s)$. Then $\|W_{t+1} - W^*\|_F^2 \leq \exp\left(-\frac{L_{t+1}}{32KD^2}\right)$.

Proof. Let $Z_t = 1$.

$$\begin{aligned}\|W_{t+1} - W^*\|_F^2 &\leq \|W_t - \eta_{t+1} \nabla \ell_t(W_t) - W^*\|_F^2 \\ &= \|W_t - W^*\|_F^2 - 2\eta_{t+1} \langle \nabla \ell_t(W_t), W_t - W^* \rangle_F + \eta_{t+1}^2 \|\nabla \ell_t(W_t)\|_F^2.\end{aligned}$$

By the separable assumption we have $\ell_t(W^*) \leq 0$. Since ℓ_t is convex, $\langle \nabla \ell_t(W_t), W_t - W^* \rangle \geq \ell_t(W_t) - \ell_t(W^*) \geq \ell_t(W_t)$. Continuing the above calculation and using Lemma 9, we get

$$\begin{aligned}\|W_{t+1} - W^*\|_F^2 &\leq \|W_t - W^*\|_F^2 - 2\eta_{t+1} \ell_t(W_t) + 8\eta_{t+1}^2 \ell_t(W_t) \\ &\leq \|W_t - W^*\|_F^2 - \frac{1}{8} \ell_t(W_t) \\ &\leq \|W_t - W^*\|_F^2 \left(1 - \frac{\ell_t(W_t)}{32KD^2}\right) \quad \text{because } \|W_t - W^*\|_F^2 \leq 4KD^2 \\ &\leq \|W_t - W^*\|_F^2 \exp\left(-\frac{\ell_t(W_t)}{32KD^2}\right)\end{aligned}$$

By induction, we can get

$$\|W_{t+1} - W^*\|_F^2 \leq KD^2 \exp\left(-\frac{L_{t+1}}{32KD^2}\right)$$

\square

Definition 11. $\|W\|_M^2 \triangleq \sum_{i=1}^K \|\mathbf{e}_i^\top W\|_M^2$.

With this definition we have $\|Wx_t\|_2^2 = \sum_{i=1}^K (\mathbf{e}_i^\top Wx_t)^2 \leq \sum_{i=1}^K \|\mathbf{e}_i^\top W\|_M^2 \|x_t\|_{M^{-1}}^2 \leq \|W\|_M^2 \|x_t\|_{M^{-1}}^2$

Lemma 12.

$$\|W_t - W^*\|_{M_t}^2 \leq (1 + L_t) K^2 D^2 \exp\left(-\frac{L_t}{32KD^2}\right) \leq 32K^3 D^4.$$

Proof. Because we assume $\|x_t\|_2^2 \leq 1$, it holds that $M_t \leq (1 + L_t)I$. Therefore $\|W_t - W^*\|_{M_t}^2 \leq (1 + L_t) \|W_t - W^*\|_I^2 = (1 + L_t) \sum_{i=1}^K \|\mathbf{e}_i^\top (W_t - W^*)\|_2^2 \leq (1 + L_t) K \|W_t - W^*\|_F^2$. By Lemma 10 this is bounded by $(1 + L_t) K^2 D^2 \exp\left(-\frac{L_t}{32KD^2}\right)$, which can further be bounded by a constant related to K and D . For example, using the property $\exp(-x) \leq \frac{1}{(1+x)^2}$ for all $x > 0$, it can be upper bounded by $(1 + L_t) K^2 D^2 \times \frac{(32KD^2)^2}{(L_t + 32KD^2)^2} \leq \frac{32^2 K^4 D^6}{32KD^2 + L_t} \leq 32K^3 D^4$. \square

Lemma 13. If $\|x_t\|_{M_t^{-1}} \leq \epsilon = \frac{1}{4D\sqrt{32K^3 D^4}}$, then $\hat{y}_t = y_t$.

Proof. By the convexity of ℓ_t ,

$$\begin{aligned} \ell_t(W_t) &\leq \ell_t(W_t) - \ell_t(W^*) \leq \langle \nabla \ell_t(W_t), W_t - W^* \rangle \\ &= \langle \nabla \Phi_t(W_t x_t) x_t^\top, W_t - W^* \rangle \\ &= \langle \nabla \Phi_t(W_t x_t), W_t x_t - W^* x_t \rangle \\ &\leq 4D \|W_t x_t - W^* x_t\|_2 \\ &\leq 4D \|W_t - W^*\|_{M_t} \|x_t\|_{M_t^{-1}} \leq 1. \end{aligned}$$

This implies $\hat{y}_t = y_t$. \square

Therefore, when we do not explore, we know W_t will predict correctly! Thus we only need to bound the number of errors occurred in exploration rounds, which is calculated by the following lemma.

Lemma 14 (The conclusion is not meaningful). $\sum_{t=1}^T \mathbf{1}[\tilde{y}_t \neq y_t] \leq \max_{t \in [T]} \left(\frac{1}{\ell_t(W_t)} \right) \times \frac{K \ln T \ln \frac{1}{\delta}}{\epsilon^2}$ with probability at least $1 - \delta$.

Proof. By the above discussion, $\sum_{t=1}^T \mathbf{1}[\tilde{y}_t \neq y_t] \leq N \triangleq \sum_{t=1}^T Z_t$, the number of exploration rounds.

$$\begin{aligned} N &= \sum_{t=1}^T \mathbf{1} \left[\|x_t\|_{M_t^{-1}} > \epsilon \right] \\ &\leq \left(K \ln \frac{1}{\delta} \right) \sum_{t=1}^T \mathbf{1} \left[\|x_t\|_{M_t^{-1}} > \epsilon \right] Z_t \quad (\text{when } \|x_t\|_{M_t^{-1}} \geq \epsilon, \tilde{y}_t = y_t \text{ with probability } \frac{1}{K}) \\ &\leq \frac{K \ln \frac{1}{\delta}}{\epsilon^2} \sum_{t=1}^T \|x_t\|_{M_t^{-1}}^2 Z_t \leq \max_{t \in [T]} \left(\frac{1}{\ell_t(W_t)} \right) \times \frac{K \ln T \ln \frac{1}{\delta}}{\epsilon^2}. \end{aligned}$$

\square

Discussion. In the calculation of Lemma 13, we can actually get $\ell_t(W_t)^2 \leq \|\nabla \Phi_t(W_t x_t)\|_2^2 \|W_t - W^*\|_{M_t}^2 \|x_t\|_{M_t^{-1}}^2$. Similar to the calculation in Lemma 10, $\|\nabla \Phi_t(W_t x_t)\|_2^2$ is bounded by constant times $\ell_t(W_t)$. So the exploration criterion could potentially become $\ell_t(W_t) \|x_t\|_{M_t^{-1}}^2 \geq \frac{1}{\epsilon^2}$, which makes Lemma 14 go through. The problem is we do not know $\ell_t(W_t)$ in general.

9 Regret Lower Bound for Explore-then-Exploit Algorithms [TODO]

10 A Regret Lower Bound for A Certain Type of Algorithms

In this section, we try to construct an error lower bound for a certain type of algorithms. This type of algorithms does not make update when it makes a wrong prediction. For simplicity, we only consider binary classification. More formally, the algorithms we consider satisfy the following assumption.

Assumption 3 (Algorithm). *Let $p_t(x)$ be the algorithm's probability of predicting class 1 (recall we consider binary classification) at round t if it receives the feature vector $x \in \mathcal{X} \subset \mathbb{R}^d$. We assume $p_t(\cdot)$ is totally determined by all previous **correct** examples. In other words, $p_t(\cdot)$ is determined by the tuple $((x_{\tau_1}, y_{\tau_1}), \dots, (x_{\tau_N}, y_{\tau_N}))$ where $1 \leq \tau_1 < \tau_2 < \dots < \tau_N < t$ are the rounds that the learner makes correct prediction.*

Assumption 4 (linearly separable with a margin). *We assume the samples are linearly separable with margin γ (i.e., any two points with different labels have distance no less than γ).*

Definition 15 (Free space). *The free space at time t is the set of points whose label is still undetermined given $(x_1, y_1), \dots, (x_{t-1}, y_{t-1})$. For example, the γ -ball centered around any already presented point is excluded from the free space. Denote the free space at time t by FS_t .*

The free space's definition simply means that at time t , the adversary can pick any point x_t in FS_t and assign the label y_t to either 1 or 2 without violating the linearly separable and the γ -margin assumption.

Below we present the Adversary's strategy of constructing (x_t, y_t) .

Algorithm 7: Adversary's strategy

```

1 Pick  $x_1$  randomly from  $\mathcal{X}$ , and let  $y_1 = 1$ .
2 for  $t = 2, \dots, T$  do
3   if  $\tilde{y}_{t-1} \neq y_{t-1}$  then
4     Let  $(x_t, y_t) = (x_{t-1}, y_{t-1})$ 
5   else if  $FS_t$  is not empty then
6     Pick  $x_t \in FS_t$ . Because of this  $x_t$ , the free space's volume is reduced. We denote the
       reduction amount by  $V_t = v(FS_{t+1}) - v(FS_t) \leq V$ . (i.e.,  $V$  is a global upper bound of  $V_t$ )
7     If  $p_t(x_t) \geq 1 - \max\left\{\sqrt{V}, \frac{1}{\sqrt{T}}\right\}$ , then label  $y_t = 2$ ; otherwise, label  $y_t = 1$ .
8   else
9     Randomly assign  $(x_t, y_t)$  with some value that does not violate the assumption.
```

Definition 16 (history). *Let \mathcal{H}_t be the history before time t : $\mathcal{H}_t = \{(x_s, y_s, \tilde{y}_s)\}_{s=1}^{t-1}$. We use $\mathbb{E}_t[\cdot]$ to denote $\mathbb{E}[\cdot | \mathcal{H}_t]$.*

Lemma 17. *If $\exists t$ such $p_t(x_t) \geq 1 - \max\left\{\sqrt{V}, \frac{1}{\sqrt{T}}\right\}$, then $\mathbb{E}_t\left[\sum_{s=t}^T \mathbf{1}[\tilde{y}_s \neq y_s]\right] \geq \Omega\left(\min\left\{\frac{1}{\sqrt{V}}, \sqrt{T}\right\}\right)$.*

Proof. By of the condition and the adversary strategy, we have $y_t = 2$. Therefore, the learner will predict the true label with probability $\leq \max\left\{\sqrt{V}, \frac{1}{\sqrt{T}}\right\}$. And note that if the learner predicts incorrectly at time t , then at time $t+1$ the feature vector remains the same ($x_{t+1} = x_t$), and the learner's probability of prediction also remains the same ($p_{t+1}(\cdot) = p_t(\cdot)$). Therefore, the expected number of mistakes before the first correct guess is (roughly) larger than $\frac{1}{\max\left\{\sqrt{V}, \frac{1}{\sqrt{T}}\right\}} = \min\left\{\frac{1}{\sqrt{V}}, \sqrt{T}\right\}$. \square

Lemma 18. *If $\forall s, p_s(x_s) \leq 1 - \max\left\{\sqrt{V}, \frac{1}{\sqrt{T}}\right\}$, then $\sum_{s=1}^T \mathbb{E}_s[\mathbf{1}[\tilde{y}_s \neq y_s]] = \Omega\left(\min\left\{\sqrt{T}, \frac{1}{\sqrt{V}}\right\}\right)$.*

Proof. By the condition and the adversary strategy, we know that the probability of error is larger than $\max \left\{ \sqrt{V}, \frac{1}{\sqrt{T}} \right\}$ at all time t before the free space is used up. Since each time the free space only reduces by V , in the first $\frac{1}{V}$ rounds (assume the total volume is 1), the free space is still all available. Therefore,

$$\sum_{s=1}^T \mathbb{E}_s[\mathbf{1}[\tilde{y}_s \neq y_s]] \geq \min \left\{ T, \frac{1}{V} \right\} \times \max \left\{ \sqrt{V}, \frac{1}{\sqrt{T}} \right\} = \min \left\{ \sqrt{T}, \frac{1}{\sqrt{V}} \right\}.$$

□

Discussion. There is a construction such that $1/V$ can be of order $\Omega \left(\left(\frac{1}{\gamma} \right)^{(d-1)/2} \right)$. [TODO] How to construct such points?

11 How Hard It Is Only Using The Weak Labels?

The halving algorithm can actually run if we can do “uniform sampling” over the version space. But it is even unknown whether we can efficiently pick a model from the version space. The problem is that we get a lot of feedback in the form of “feature x_t does not belong to class \tilde{y}_t ”, which we don’t know how to use.

The following is just an attempt (not successful but might be interesting...) to say that it might be not easy to figure out a model if the learner is only presented with this kind of “error message”.

The problem is formulated as follows. Given N points in a row, each one with a class $c_i \in [K]$, $\forall i \in [N]$. We call these N points *separable* if the following statement holds:

$$\text{If } c_i = c_j \text{ for some } i \leq j, \text{ then } c_i = c_{i+1} = \dots = c_j.$$

For example, if $N = 5, K = 3$, then $(c_1, c_2, c_3, c_4, c_5) = (3, 3, 1, 1, 1)$ is separable, but $(c_1, c_2, c_3, c_4, c_5) = (2, 1, 2, 2, 2)$ is not.

Now you have N conditions, in which the i -th condition only says something like “ $c_i \neq k$ ” for some k .

(1) Can you efficiently decide whether there exists an assignment of (c_1, \dots, c_N) such that these N points are separable and satisfy all the conditions?

(2) If it is guaranteed that there are separable solutions, can you efficiently find one of them?

(Efficient: the complexity is polynomial in N and K)

Example 1.

$N = 5, K = 3$:

$$c_1 \neq 1$$

$$c_2 \neq 2$$

$$c_3 \neq 3$$

$$c_4 \neq 1$$

$$c_5 \neq 2$$

$\Rightarrow (c_1, c_2, c_3, c_4, c_5) = (2, 1, 1, 3, 3)$ or $(3, 3, 2, 2, 1)$ are separable solutions.

Example 2.

$N = 7, K = 3$:

$$c_1 \neq 1$$

$$c_2 \neq 2$$

$$c_3 \neq 3$$

$$c_4 \neq 1$$

$$c_5 \neq 2$$

$$c_6 \neq 3$$

$$c_7 \neq 1$$

\Rightarrow There is no separable solution.

It turns out this specific 1-dimensional problem is equivalent to identify a **missing permutation** of $[K]$ as a subsequence in the weak-label sequence. In Example 1, the existing permutations are $(1, 2, 3)$, $(1, 3, 2)$, $(2, 3, 1)$, $(3, 1, 2)$, the missing ones are $(2, 1, 3)$ and $(3, 2, 1)$. So the solutions can be $(2, 1, 3)$ or $(3, 2, 1)$ (with some repetition). In Example 2, all permutations appear in the weak-label sequence, so there is no solution.

We can prove this equivalence considering two directions:

- (1) If there is a solution whose class labels follow a permutation, then that permutation cannot be a subsequence of the weak-label sequence.
- (2) If there is a missing permutation in the weak-label sequence, then there is a class label assignment that follows this permutation.

[TODO] Proof

12 Biased Halving: Trading Error with Complexity

Algorithm 8: Banditron

```

1 Define:  $\Omega = \{W \in \mathbb{R}^{Kd} : \|\mathbf{e}_i^\top W\|_2 \leq D\}$ .
2 For a set  $S$  of  $W$ 's,  $S(i|x)$  is the subset of  $S$  that outputs class  $i$  given feature vector  $x$ , i.e.,
    $S(i|x) = \{W \in S : (Wx)_i \geq (Wx)_j \forall j\}$ 
3  $|S|$  denotes the volume of  $S$ .
4 parameter:  $\alpha \in (0, 1)$ 
5  $\Omega_1 = \Omega$ .
6 for  $t = 1, \dots, T$  do
7   if  $\arg\max_i \frac{|\Omega_t(i|x_t)|}{|\Omega_t|} \geq 1 - \alpha$  then
8     Let  $\tilde{y}_t = i$ .
9     if  $\tilde{y}_t \neq y_t$  then
10       $\Omega_{t+1} = \Omega_t \setminus \Omega_t(\tilde{y}_t|x_t)$ .
11   else
12     Let  $\tilde{y}_t \sim \text{unif}([K])$ .
13     if  $\tilde{y}_t = y_t$  then
14       $\Omega_{t+1} = \Omega_t(\tilde{y}_t|x_t)$ .

```

Conjecture(should be true): If the volume of Ω_t becomes smaller than $\frac{|\Omega|}{N}$, then the algorithm won't make any error anymore. N should be in the order of $\Theta(\frac{1}{\gamma^{Kd}})$.

Rough analysis:

Each time the algorithm makes an error in Line 7, the volume becomes α times the original volume. So the algorithm will not make more than $\frac{\ln N}{\ln \frac{1}{\alpha}}$ mistakes in this case.

In the case of Line 10, $K \ln \frac{1}{\delta}$ errors will accompany with a $(1 - \alpha)$ -factor shrinkage in the volume.

Therefore, the number of errors occurred in this case is upper bounded by $\frac{K \ln \frac{1}{\delta} \ln N}{\ln \frac{1}{1-\alpha}} \leq \frac{K \ln \frac{1}{\delta} \ln N}{\alpha}$.

Now we discuss about the complexity. The main issue is how to maintain Ω_t . Each time the algorithm enters Line 10, Ω_t becomes more and more fragmented. But if Ω_t can be maintained with M convex cones, then Ω_{t+1} can be maintained with $(K - 1)M \leq KM$ convex cones. And we assume each cone's volume can be computed in $\text{poly}(T)$ time. Each time the algorithm enters Line 14, the number of convex cones does not increase.

By the above discussion, there will be no more than $K \frac{\ln N}{\ln \frac{1}{\alpha}}$ convex cones to maintain. And the error bound is in the order of $\frac{K \ln \frac{1}{\delta} \ln N}{\alpha}$ for some $\alpha < \frac{1}{2}$. Let's try to balance the number of errors and computational complexity. Let

$$\begin{aligned}
K \frac{\ln N}{\ln \frac{1}{\alpha}} &\approx \frac{K \ln \frac{1}{\delta} \ln N}{\alpha} \\
\Rightarrow \frac{\ln N}{\ln \frac{1}{\alpha}} \ln K &\approx \ln \left(K \ln \frac{1}{\delta} \ln N \right) + \ln \frac{1}{\alpha}
\end{aligned}$$

$$\Rightarrow \text{pick } \ln \frac{1}{\alpha} = \sqrt{\ln N}.$$

Thus the computational complexity is in the order of $K^{\sqrt{\ln N}} \times \text{poly}(T) = K^{\sqrt{Kd \ln \frac{1}{\delta}}}$. The error bound is $\mathcal{O}\left(e^{\sqrt{Kd \ln \frac{1}{\delta}}} K^2 d \ln \frac{1}{\delta} \ln \frac{1}{\gamma}\right)$.

Another viewpoint: let $\frac{1}{\alpha} = K^\beta$, then the complexity is $\left(\frac{1}{\gamma}\right)^{\frac{Kd}{\beta}} \times \text{poly}(T)$ and the error bound is $K^{\beta+1} \ln \frac{1}{\delta} \ln N$.

13 Continuous EXP4 with Uniform Exploration

Algorithm 9: Banditron

- 1 **Parameters:** feasible set $\Omega \subset \mathbb{R}^{K \times d}$
 - 2 **Definitions:** $\ell_t(W) \triangleq [1 - (Wx_t)_{y_t} + \max_{r \in [K]} (Wx_t)_r]_+$ (hinge loss)
 - 3 **for** $t = 1, \dots, T$ **do**
 - 4 Receive $x_t \in \mathbb{R}^d$.
 - 5 Define

$$q_t(W) = \frac{\exp(-\alpha \sum_{s=1}^{t-1} \hat{\ell}_s(W))}{\int_{U \in \Omega} \exp(-\alpha \sum_{s=1}^{t-1} \hat{\ell}_s(U)) dU}, \quad \forall W \in \Omega,$$

where $\hat{\ell}_s(W) = \mathbf{1}[\tilde{y}_s = y_s] \left(\frac{\mathbf{1}[\tilde{y}_s = y_s] \ell_s(W)}{1 - \gamma + \frac{\gamma}{K}} + \frac{\mathbf{1}[\tilde{y}_s \neq y_s] \ell_s(W)}{\frac{\gamma}{K}} \right)$.
 - 6 Sample $W_t \sim q_t$, and let $\hat{y}_t = \arg\max_{r \in [K]} (W_t x_t)_r$.
 - 7 Let $\tilde{y}_t = \hat{y}_t$ with probability $1 - \gamma$, and $\tilde{y}_t \sim \text{unif}([K])$ with probability γ .
-

Lemma 19. $\mathbb{E}_{\tilde{y}_t}[\hat{\ell}_t(W)] = \ell_t(W)$ for all W .

Proof.

$$\begin{aligned} \mathbb{E}_{\tilde{y}_t}[\hat{\ell}_t(W)] &= \mathbb{E}_{\tilde{y}_t} \left[\mathbf{1}[\hat{y}_t = y_s] \frac{\mathbf{1}[\tilde{y}_s = y_s] \ell_t(W)}{1 - \gamma + \frac{\gamma}{K}} + \mathbf{1}[\hat{y}_t \neq y_s] \frac{\mathbf{1}[\tilde{y}_s = y_s] \ell_t(W)}{\frac{\gamma}{K}} \right] \\ &= \mathbf{1}[\hat{y}_t = y_s] \ell_t(W) + \mathbf{1}[\hat{y}_t \neq y_s] \ell_t(W) = \ell_t(W). \end{aligned}$$

□

Plugging these lemmas in the previous hedge bound, we can get

$$\begin{aligned} & \mathbb{E} \left[\sum_{t=1}^T \ell_t(W_t) \right] \\ &= \mathbb{E} \left[\sum_{t=1}^T \int_{W \in \Omega} q_t(W) \ell_t(W) dW \right] \\ &= \mathbb{E} \left[\sum_{t=1}^T \mathbb{E}_{\tilde{y}_t} \left[\int_{W \in \Omega} q_t(W) \hat{\ell}_t(W) dW \right] \right] \\ &\leq \mathbb{E} \left[\sum_{t=1}^T \mathbb{E}_{\tilde{y}_t} [\hat{\ell}_t(W^*)] + \frac{\mathbf{Ent}(q_1 || \delta(W^*))}{\alpha} + \alpha \sum_{t=1}^T \mathbb{E}_{\tilde{y}_t} \left[\int_{W \in \Omega} q_t(W) \hat{\ell}_t(W)^2 dW \right] \right] \\ &\leq \mathbb{E} \left[\sum_{t=1}^T \mathbb{E}_{\tilde{y}_t} [\hat{\ell}_t(W^*)] + \frac{\mathbf{Ent}(q_1 || \delta(W^*))}{\alpha} + \frac{2K\alpha}{\gamma} \sum_{t=1}^T \mathbb{E}_{\tilde{y}_t} \left[\int_{W \in \Omega} q_t(W) \hat{\ell}_t(W) dW \right] \right] \end{aligned}$$

... to bound the regret, it would be something like bounding $\frac{1}{1 - \frac{K\alpha}{\gamma}} \left(\frac{1}{\alpha} + \frac{K\alpha}{\gamma} L^* + \gamma T \right)$, which gives $(L^* T)^{1/3} + \sqrt{T}$ regret bound.

Discussion. We can change $\ell_t(\cdot)$ to any reasonable convex loss (e.g., logsitic loss or second-order loss).

References

- [1] Alekh Agarwal. Selective sampling algorithms for cost-sensitive multiclass prediction. In *International Conference on Machine Learning*, pages 1220–1228, 2013.
- [2] Alina Beygelzimer, Francesco Orabona, and Chicheng Zhang. Efficient online bandit multiclass learning with $\tilde{O}(\sqrt{T})$ regret. In *International Conference on Machine Learning*, 2017.
- [3] Koby Crammer and Claudio Gentile. Multiclass classification with bandit feedback using adaptive regularization. *Machine learning*, 90(3):347–383, 2013.
- [4] Sarah Filippi, Olivier Cappé, Aurélien Garivier, and Csaba Szepesvári. Parametric bandits: The generalized linear case. In *Advances in Neural Information Processing Systems*, pages 586–594, 2010.
- [5] D. J. Foster and A. Krishnamurthy. Contextual bandits with surrogate losses: Margin bounds and efficient algorithms. *ArXiv e-prints*, June 2018.
- [6] Dylan J Foster, Satyen Kale, Haipeng Luo, Mehryar Mohri, and Karthik Sridharan. Logistic regression: The importance of being improper. *Proceedings of Machine Learning Research*, 75:1–42, 2018.
- [7] Elad Hazan and Satyen Kale. Newtron: an efficient bandit algorithm for online multiclass prediction. In *Advances in neural information processing systems*, pages 891–899, 2011.
- [8] Sham M Kakade, Shai Shalev-Shwartz, and Ambuj Tewari. Efficient bandit algorithms for online multiclass prediction. In *Proceedings of the 25th International Conference on Machine Learning*, 2008.
- [9] Adam R Klivans and Rocco A Servedio. Learning intersections of halfspaces with a margin. In *International Conference on Computational Learning Theory*, 2004.
- [10] Loucas Pillaud-Vivien, Alessandro Rudi, and Francis Bach. Exponential convergence of testing error for stochastic gradient methods. 2017.